

INFO20003 Tutorial 5

starting ~ 2.20 pm

Today's tutorial

- Review of Relational Algebra (RA)
- RA & SQL questions
 - group work

Github Link

<https://github.com/Prashansa-Singh/INFO20003-Sem1-2022>

Relational Algebra (RA) Review:

RA

- Consists of a collection of operators
- These take instances of a relation (tables) as their inputs
- Return an instance of a relation as output
- Can be unary or binary operators
 - unary: applied on a single relation
 - binary: requires two relations as input

Five fundamental basic operations of RA

- Selection
- Projection
- Cross Product
- Set Difference
- Set Union

- These five basic operations can be used to form compound operations (Intersection, Natural Join and Condition Join)

Removal operations

- Selection
- Projection
- These remove components from a relation
- Selection: removes rows
- Projection: removes columns

Projection

$$\pi_{A_1, A_2, \dots, A_n}(R)$$

relation

Attributes / columns of relation which are projected (kept)

- Creates a new relation with a subset of the columns
- All tuples from original relation are kept, but we only store projected columns

- If one of the columns which is removed is part of the PK
 - o You may end up with duplicate rows in your resulting relation
 - o Projection operator in relational algebra automatically removes duplicates rows

- Example:

- o Person table

FirstName	LastName	Phone	Email
Jon	Snow	0551-999-210	knowsnothing@hotmail.com
Daenerys	Targaryen	0569-988-112	bendtheknee@gmail.com
Jamie	Lannister	0531-987-654	handsfree@gmail.com
Night	King	0566-123-456	killerstare@gmail.com

The expression $\pi_{\text{FirstName, LastName}}(\text{Person})$ will result in:

- Result of projection example:

FirstName	LastName
Jon	Snow
Daenerys	Targaryen
Jamie	Lannister
Night	King

Selection

$\sigma_{\text{condition}}(R)$

- R is a relation
- Condition is used to filter rows
- Creates a new relation which only includes the rows for which the condition is true

- Example

- o Person table

FirstName	LastName	Phone	Email
Jon	Snow	0551-999-210	knowsnothing@hotmail.com
Daenerys	Targaryen	0569-988-112	bendtheknee@gmail.com
Jamie	Lannister	0531-987-654	handsfree@gmail.com
Night	King	0566-123-456	killerstare@gmail.com

- o $\sigma_{\text{FirstName} = \text{'Jon'} \vee \text{LastName} = \text{'King'}}(\text{Person})$

$\vee = \text{or}$

$\wedge = \text{And}$

▪ Result of Selection example

FirstName	LastName	Phone	Email
Jon	Snow	0551-999-210	knowsnothing@hotmail.com
Night	King	0566-123-456	killerstare@gmail.com

We can combine the two operations in one expression as:

$\pi_{\text{FirstName, LastName}} (\sigma_{\text{FirstName} = \text{'Jon'} \vee \text{LastName} = \text{'King'}} (\text{Person}))$

FirstName	LastName
Jon	Snow
Night	King

- **Set operations**

- Cross Product
- Set Difference
- Set Union
- Both Union and Set Difference require input relations to be **union-compatible**
 - Same number of attributes in same order
 - Corresponding attributes have same data type

Union

- R U S
- Result is every row which is either in R **or** S
 - (Duplicate rows are not stored)

- Example:

GoodGuys

BadGuys

FirstName	LastName	FirstName	LastName
Jon	Snow	Cersei	Lannister
Daenerys	Targaryen	Night	King

GoodGuys U BadGuys will result in:

GoodGuys \cup BadGuys will result in:



FirstName	LastName
Jon	Snow
Daenerys	Targaryen
Cersei	Lannister
Night	King

Set Difference

- **R - S**

- Result is every row which is **in R but not S**

- S-R (order matters in set difference - this returns a different result to R - S)

- Example:

FirstName	LastName
Jon	Snow
Daenerys	Targaryen
Jamie	Lannister
Night	King

FirstName	LastName
Night	King
Arya	Stark
Cersei	Lannister
Daenerys	Targaryen

RandomCombo1 - RandomCombo2 will result in:

FirstName	LastName
Jon	Snow
Jamie	Lannister

Cross Product

- R x S

- Each row in R is paired with each row in S

- Resulting relation has all attributes from both relations

- o (If some attributes have same name, rename them using the renaming operator ρ 'rho' - will learn later in tute)

- Example

Person			Weapon	
FirstName	LastName	Email	Weapon	Metal
Jon	Snow	knowsnothing@hotmail.com	Sword	Valyrian steel
Night	King	killerstare@gmail.com	Dagger	Dragon glass

Person

FirstName	LastName	Email
Jon	Snow	knowsnothing@hotmail.com
Night	King	killerstare@gmail.com

Weapon

Weapon	Metal
Sword	Valyrian steel
Dagger	Dragon glass

Person \times Weapon will result in:

$$2 \times 2 = 4$$

- How many rows will be in the result?

FirstName	LastName	Email	Weapon	Metal
Jon	Snow	knowsnothing@hotmail.com	Sword	Valyrian steel
Jon	Snow	knowsnothing@hotmail.com	Dagger	Dragon glass
Night	King	killerstare@gmail.com	Sword	Valyrian steel
Night	King	killerstare@gmail.com	Dagger	Dragon glass

-

Compound Operations

- Can all be expressed in terms of basic operators

Intersection

- Is a set operator
- Result is a relation containing all rows which are present in **both** relations
- 2 relations should be union compatible

- Intersection in terms of basic operations:

$$R \cap S = R - (R - S)$$

- Example

FirstName	LastName
Jon	Snow
Daenerys	Targaryen
Jamie	Lannister
Night	King

FirstName	LastName
Night	King
Arya	Stark
Cersei	Lannister
Daenerys	Targaryen

RandomCombo1 \cap RandomCombo2 will result in:

FirstName	LastName
Daenerys	Targaryen
Night	King

Natural Join

- $R \bowtie S$
- First, we find attributes common to both relations (have the same name)
 - o Equality is implicit
- Then create a new relation, pairing each tuple from R and S where common attributes are equal
- Joins are compound operators
 - o usually they are made up of the basic operators
 - cross product
 - selection
 - and projection
- Can break down a NJ into these steps:
 - 1) Compute $R \times S$
 - 2) Select rows where attributes that appear in both relations have equal values
 - 3) Project all unique attributes and keep one copy of the common attributes

For example, here is a natural join between the Person and WeaponOwner relations:

Person			WeaponOwner		
FirstName	LastName	Email	Weapon	LastName	Metal
Jon	Snow	knowsnothing@hotmail.com	Sword	Snow	Valyrian steel
Daenerys	Targaryen	bendtheknee@gmail.com	Dagger	Lannister	Dragon glass
Tyrion	Lannister	idrinkandiknow@gmail.com			
Night	King	killerstare@gmail.com			

- Result:
 - o Method 1)
 - Find attributes common to both relations
 - Then create a new relation, pairing each tuple from R and S where common attributes are equal

- o Method 2)
 - Break down into steps using basic operations

Person \times Weapon (intermediate result):

FirstName	LastName	Email	Weapon	LastName	Metal
Jon	Snow	knowsnothing@hotmail.com	Sword	Snow	Valyrian steel
Jon	Snow	knowsnothing@hotmail.com	Dagger	Lannister	Dragon glass
Daenerys	Targaryen	bendtheknee@gmail.com	Sword	Snow	Valyrian steel
Daenerys	Targaryen	bendtheknee@gmail.com	Dagger	Lannister	Dragon glass
Tyrion	Lannister	idrinkandiknow@gmail.com	Sword	Snow	Valyrian steel
Tyrion	Lannister	idrinkandiknow@gmail.com	Dagger	Lannister	Dragon glass
Night	King	killerstare@gmail.com	Sword	Snow	Valyrian steel
Night	King	killerstare@gmail.com	Dagger	Lannister	Dragon glass

Person \bowtie Weapon will result in:

Night	King	killerstare@gmail.com	Dagger	Lannister	Dragon glass
-------	------	-----------------------	--------	-----------	--------------

Person \bowtie Weapon will result in:

FirstName	LastName	Email	Weapon	Metal
Jon	Snow	knowsnothing@hotmail.com	Sword	Valyrian steel
Tyrion	Lannister	idrinkandiknow@gmail.com	Dagger	Dragon glass

Notice LastName attribute only appears once

Condition Join

- Also called Theta/Inner Join

- $R \bowtie_C S$

- This joins rows from R and S such that the condition C is true
- Usually, C is of type $A = B$ (this is an equi-join, which is a type of condition join)
- Condition joins can also involves $!=$, $<$, $>$ etc e.g. Attribute1 $!=$ 5

- Condition Join in terms of basic operations:

- $R \bowtie_C S = \sigma_C(R \times S)$

Example:

Person			WeaponOwner		
FirstName	LastName	Email	Weapon	Name	Metal
Jon	Snow	knowsnothing@hotmail.com	Sword	Snow	Valyrian steel
Daenerys	Targaryen	bendtheknee@gmail.com	Dagger	Lannister	Dragon glass
Tyrion	Lannister	idrinkandiknow@gmail.com			
Night	King	killerstare@gmail.com			

Person $\bowtie_{\text{LastName} = \text{Name}}$ Weapon will result in:

- Why isn't this considered a natural join?

- This is very similar to a Natural Join but the difference is you can't use natural join because attribute names are different.
- Natural Joins
 - Equality is implicit. (You don't need to specify which attributes you're joining on)
 - Attributes **must have same name** for NJ to work
 - NJ has implicit projection - only 1 copy of attribute is kept
- Equi-joins
 - Equality is explicit, can make up any equality condition you want
 - If both attributes have same name, you need to explicitly project to keep one

Person × Weapon (intermediate result):

FirstName	LastName	Email	Weapon	Name	Metal
Jon	Snow	knowsnothing@hotmail.com	Sword	Snow	Valyrian steel
Jon	Snow	knowsnothing@hotmail.com	Dagger	Lannister	Dragon glass
Daenerys	Targaryen	bendtheknee@gmail.com	Sword	Snow	Valyrian steel
Daenerys	Targaryen	bendtheknee@gmail.com	Dagger	Lannister	Dragon glass
Tyion	Lannister	idrinkandiknow@gmail.com	Sword	Snow	Valyrian steel
Tyion	Lannister	idrinkandiknow@gmail.com	Dagger	Lannister	Dragon glass
Night	King	killerstare@gmail.com	Sword	Snow	Valyrian steel
Night	King	killerstare@gmail.com	Dagger	Lannister	Dragon glass

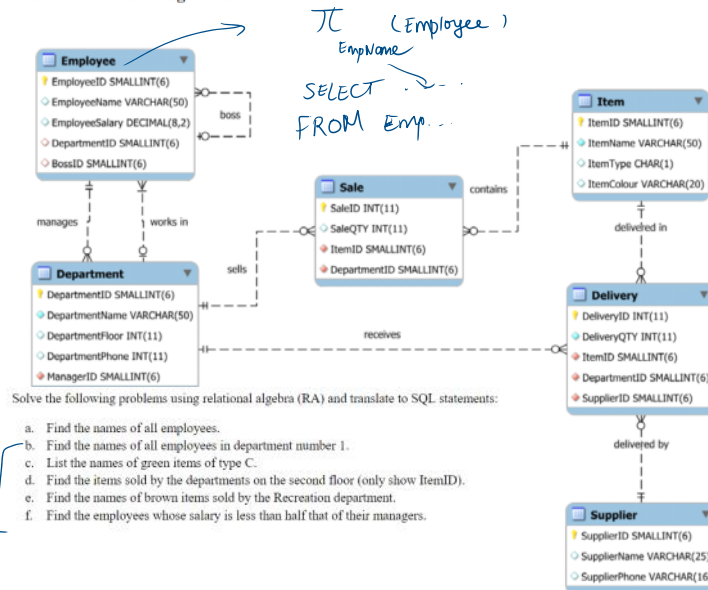
Person ⋈_{LastName = Name} Weapon will result in:

FirstName	LastName	Email	Weapon	Name	Metal
Jon	Snow	knowsnothing@hotmail.com	Sword	Snow	Valyrian steel
Tyion	Lannister	idrinkandiknow@gmail.com	Dagger	Lannister	Dragon glass

Notice both the Name and LastName attributes appear in result

Q2


Consider the following schema:



Solve the following problems using relational algebra (RA) and translate to SQL statements:

- Find the names of all employees.
- Find the names of all employees in department number 1.
- List the names of green items of type C.
- Find the items sold by the departments on the second floor (only show ItemID).
- Find the names of brown items sold by the Recreation department.
- Find the employees whose salary is less than half that of their managers.

How to write RA statements:



- Write and read from the inside out (start with innermost bracket and work your way out)

- STEP 1

○ TABLES/JOINS

- First figure out all the relations you need to use to access the attributes you want. These go inside the innermost bracket
 - Figure out how these must get joined
 - e.g. natural join, equijoin, condition join etc
 - You might just need 1 table and so don't need any joins

- STEP 2

○ SELECTION

- Then after we have the table / result of joined tables we want, we do **selection on the rows** of this result. This is second inner most bracket.
- Figure out which conditions we are doing selection on (horizontal filtering)
 - These are the **conditions we must ensure all of our returned rows in our final output result fulfil**
 - e.g. for part
 - ◆ b) dept number is 1
 - ◆ c) item colour is green and item type is c
 - ◆ d) dept floor is 2
 - ◆ e) item colour is brown and dept name is Recreation
 - ◆ f) emp.salary is $< 1/2$ of manager salary

- STEP 3

○ PROJECTION

- Figure out the columns/attributes you actually want in your final result
 - Do we just want to see the employee names? or names of items? or ItemID values? etc
 - This the very last step (performed after joins and selection) and will keep only these columns in your output results
 - Done in outermost brackets

Translating from RA to SQL

- SELECT {Columns}

- This statement is actually the last to get executed
- It takes the result after the tables have been joined (FROM step) and selection on the rows of this result has occurred (WHERE step), and only keeps the columns specified
- This is equivalent to the **Projection step** (list the columns you want to keep)

- FROM RelationA NATURAL JOIN RelationB NATURAL JOIN RelationC ...

- A INNER JOIN B ON LastName = Name

- This is equivalent to the **Tables/Joins step** and is executed first

○ This returns a table of records which is the output of all these joins

- **WHERE {Conditions}**

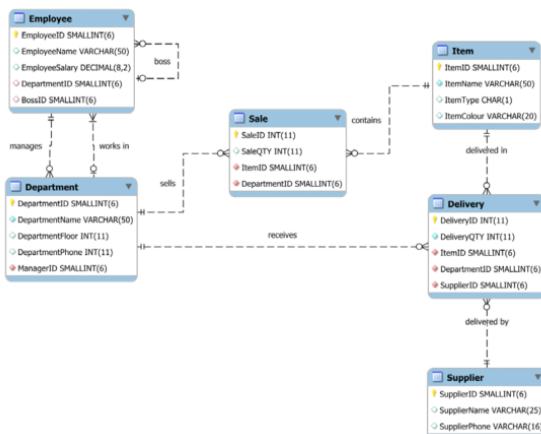
○ This step is performed after the FROM step and before SELECT

○ This is equivalent to the **Selection step**

○ In this step, we use the output result we got in the FROM statement and only keep the rows in this result which fulfil the conditions specified

- Beware: You can have more complicated SQL statements with more than these 3 keywords (SELECT, FROM, WHERE). This is just a template for a basic SQL statement.

a. Find the names of all employees.



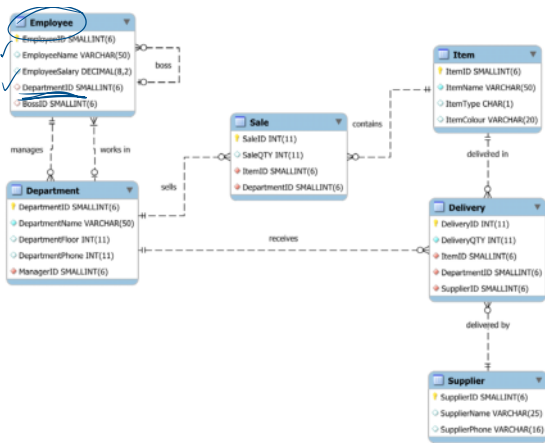
a
SELECT EmpName
FROM Employee;

Semi-colon (;) marks end of query

• SELECT in SQL means projection in RA

b. Find the names of all employees in department number 1.



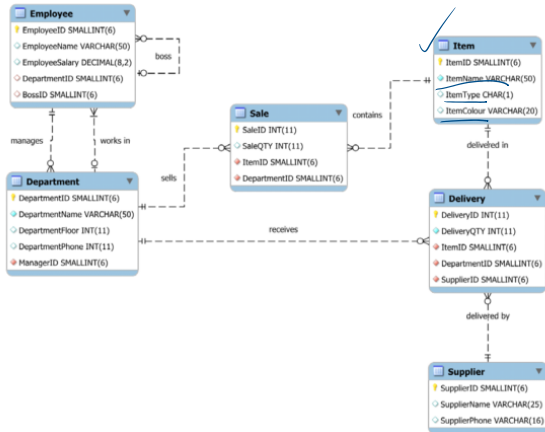


b

• Best to do **projection after selection!**

- What happens if you do it the other way around?
- If the attributes used in selection condition are not projected, then selection does not work!

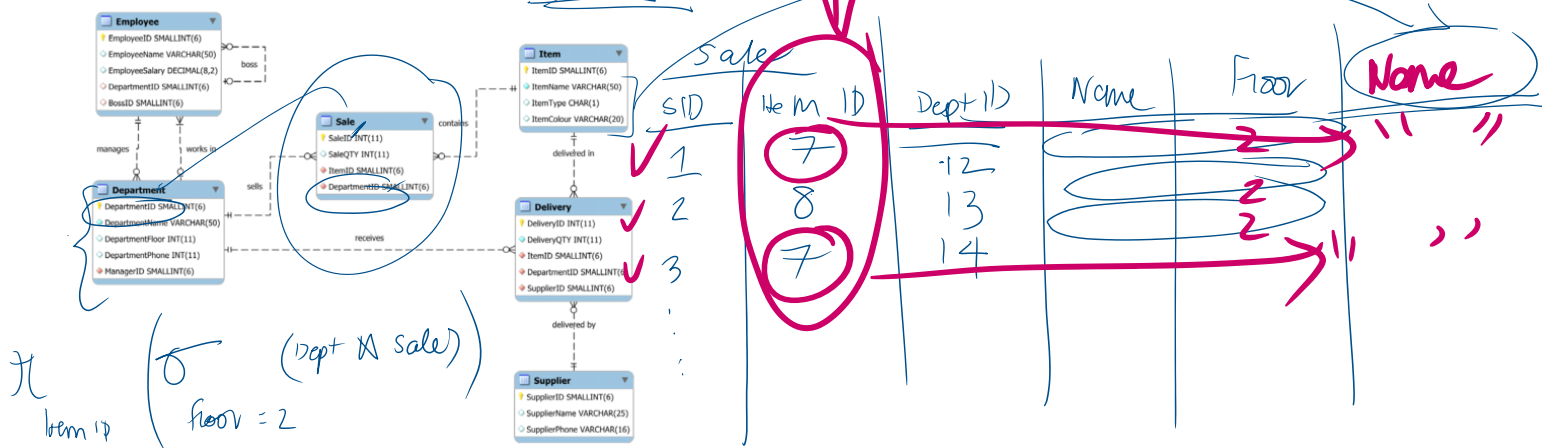
c. List the names of green items of type C.



C

- Should use single quotes for strings e.g. 'Green'

d. Find the items sold by the departments on the second floor (only show ItemID)



d

```
SELECT DISTINCT ItemID
FROM Department NATURAL JOIN Sale
WHERE DepartmentFloor = 2;
```

- Why do we use DISTINCT in SQL?

- This returns only the different values
- Inside a table, columns can contain many duplicate values
- We only care about the distinct values
- MySQL doesn't discard duplicates, so we must use DISTINCT to do so
- DISTINCT
 - Sale captures each individual sale
 - A specific item can be sold many times
 - (Only SaleID is the PK)

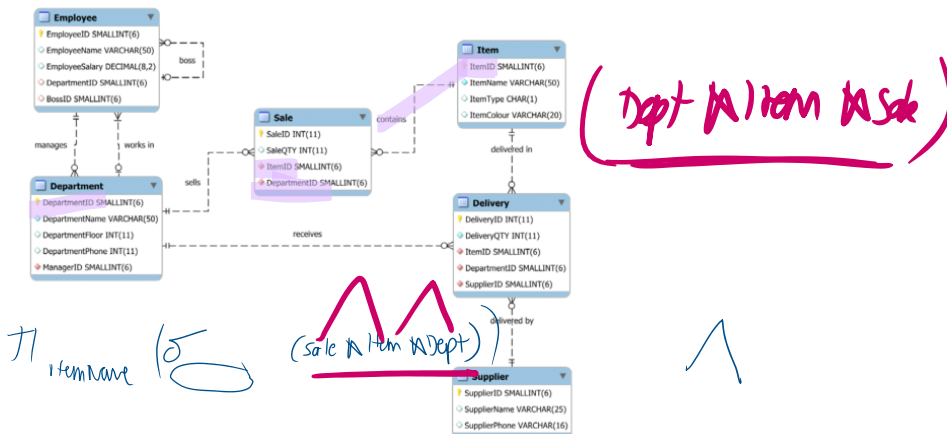
- When you natural join Sale and Department
 - The common attribute is DepartmentID
 - The resulting schema contains all the other columns
 - (Basically you're now storing the department information for each sale)
- So in the column ItemID, there may be lots of duplicate values because the same item has been sold many items
 - We don't want our list of ItemIDs to contain duplicates, we only care about the unique/distinct items sold

Sale NJ Dept

SaleID	ItemID	SaleQTY	DeptID	DeptName
1	1				
2	1				
3	1				
4	1				
...					

- In RA,
 - Projection (automatically) removes duplicate statements
 - so no need to worry about doing DISTINCT somehow in relational algebra
 - DISTINCT is only for SQL

e. Find the names of brown items sold by the Recreation department.

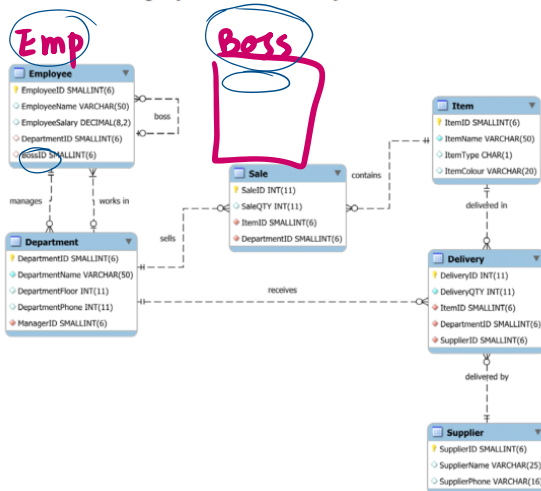


e

```
SELECT DISTINCT ItemName
FROM SALE NJ ITEM NJ DEPARTMENT
WHERE DeptName = 'Recreation' AND ItemColour = 'Brown';
```

- You should also use DISTINCT in the SQL statement here
 - Natural Join gives us the Sale Table, but also storing information on the item and department of each sale
 - Items can appear many times like in d)
 - We only want the distinct item names

f. Find the employees whose salary is less than half that of their managers.



f

- Unary relationship/join occurs
- Many ways to do this - all involve some sort of renaming of tables

Solution for f)

f. Find the employees whose salary is less than half that of their managers.

Relational Algebra: Note: The RA notation for unary joins is not agreed upon

Here are two examples using the rename (ρ) operator:

$\rho(\text{Emp}(\text{EmployeeName} \rightarrow \text{EmpName}, \text{EmployeeSalary} \rightarrow \text{EmpSalary}, \text{BossID} \rightarrow \text{EmpBossID}), \text{Employee})$ $\rho(\text{Boss}(\text{EmployeeID} \rightarrow \text{BossEmployeeID}, \text{EmployeeSalary} \rightarrow \text{BossSalary}), \text{Employee})$ $\pi_{\text{EmpName}} (\sigma_{\text{EmpSalary} < (\text{BossSalary} / 2)} (\text{Emp} \bowtie_{\text{EmpBossID} = \text{BossEmployeeID}} \text{Boss}))$	$\rho(\text{Boss}(\text{BossID} \rightarrow \text{BossBossID}, \text{EmployeeID} \rightarrow \text{BossID}, \text{Salary} \rightarrow \text{BossSalary}, \text{Name} \rightarrow \text{BossName}), \text{Employee})$ $\pi_{\text{EmployeeName}} (\sigma_{\text{EmployeeSalary} < (\text{BossSalary} / 2)} (\text{Employee} \bowtie \text{Boss}))$
---	---

- LHS: renames both tables
- RHS: renames only one table
 - becomes a natural join as same attribute name
- For both LHS and RHS, you are not using notation such as table.attributeName like you are doing the the below solution using SQL notation.
 - This is because all the relevant attribute names in the two tables have been changed so they are unique.
 - So you only need to say EmployeeSalary in RHS, don't need to say Employee.EmployeeSalary because that attribute name only appears in the Employee table.

Another way:

Or you could use an SQL-like notation:

Emp := Employee
 Boss := Employee

$$\pi_{\text{Emp.EmployeeName}} (\sigma_{\text{Emp.EmployeeSalary} < (\text{Boss.EmployeeSalary} / 2)} (\text{Emp} \bowtie_{\text{Emp.BossID} = \text{Boss.EmployeeID}} \text{Boss}))$$

SQL: **SELECT** Emp.EmployeeName
FROM Employee AS Emp
INNER JOIN Employee AS Boss
ON Emp.BossID = Boss.EmployeeID
WHERE Emp.EmployeeSalary < (Boss.EmployeeSalary / 2);

Inner join in SQL is equi-join in RA

Solutions

- a. Find the names of all employees.

Relational Algebra: $\pi_{\text{EmployeeName}} (\text{Employee})$

SQL: **SELECT** EmployeeName
FROM Employee;

- b. Find the names of all employees in department number 1.

Relational Algebra: $\pi_{\text{EmployeeName}} (\sigma_{\text{DepartmentID} = 1} (\text{Employee}))$

SQL: **SELECT** EmployeeName
FROM Employee
WHERE DepartmentID = 1;

- c. List the names of green items of type C.

Relational Algebra: $\pi_{\text{ItemName}} (\sigma_{\text{ItemColour} = \text{'Green'} \wedge \text{ItemType} = \text{'C'}} (\text{Item}))$

SQL: **SELECT** ItemName
FROM Item
WHERE ItemType = 'C' AND ItemColour = 'Green';

- We don't need distinct here in (c)
 - All ItemNames are expected to be distinct since ItemID is the PK
 - But can add it in and it won't change the result

- d. Find the items sold by the departments on the second floor (only show ItemID).

Relational Algebra: $\pi_{\text{ItemID}} (\sigma_{\text{DepartmentFloor} = 2} (\text{Sale} \bowtie \text{Department}))$

SQL: **SELECT DISTINCT** ItemID
FROM Sale **NATURAL JOIN** Department
WHERE DepartmentFloor = 2;

- e. Find the names of brown items sold by the Recreation department.

Relational Algebra: $\pi_{\text{ItemName}} (\sigma_{\text{DepartmentName} = \text{'Recreation'} \wedge \text{ItemColour} = \text{'Brown'}} (\text{Item} \bowtie \text{Sale} \bowtie \text{Department}))$

SQL: **SELECT** ItemName
FROM Item **NATURAL JOIN** Sale **NATURAL JOIN** Department
WHERE DepartmentName = 'Recreation'
AND ItemColour = 'Brown';

** I would use DISTINCT in part e) as well

** **SELECT DISTINCT** ItemName

- f. Find the employees whose salary is less than half that of their managers.

Relational Algebra: *Note: The RA notation for unary joins is not agreed upon*

Here are two examples using the rename (ρ) operator:

$\rho(\text{Emp}(\text{EmployeeName} \rightarrow \text{EmpName}, \text{EmployeeSalary} \rightarrow \text{EmpSalary}, \text{BossID} \rightarrow \text{EmpBossID}), \text{Employee})$ $\rho(\text{Boss}(\text{EmployeeID} \rightarrow \text{BossEmployeeID}, \text{EmployeeSalary} \rightarrow \text{BossSalary}), \text{Employee})$ $\pi_{\text{EmpName}} (\sigma_{\text{EmpSalary} < (\text{BossSalary} / 2)} (\text{Emp} \bowtie_{\text{EmpBossID} = \text{BossEmployeeID}} \text{Boss}))$	$\rho(\text{Boss}(\text{BossID} \rightarrow \text{BossBossId}, \text{EmployeeID} \rightarrow \text{BossID}, \text{Salary} \rightarrow \text{BossSalary}, \text{Name} \rightarrow \text{BossName}), \text{Employee})$ $\pi_{\text{EmployeeName}} (\sigma_{\text{EmployeeSalary} < (\text{BossSalary} / 2)} (\text{Employee} \bowtie \text{Boss}))$
---	--

Or you could use an SQL-like notation:

Emp := Employee

Boss := Employee

$\pi_{\text{Emp.EmployeeName}} (\sigma_{\text{Emp.EmployeeSalary} < (\text{Boss.EmployeeSalary} / 2)} (\text{Emp} \bowtie_{\text{Emp.BossID} = \text{Boss.EmployeeID}} \text{Boss}))$

SQL: `SELECT Emp.EmployeeName
FROM Employee AS Emp
INNER JOIN Employee AS Boss
ON Emp.BossID = Boss.EmployeeID
WHERE Emp.EmployeeSalary < (Boss.EmployeeSalary / 2);`